


Announcements

- Homework 1: Search
 - On edX - online, instant grading, submit as often as you like.
- Project 1: Search
 - Start early and ask questions. It's longer than most!

1

CS 188: Artificial Intelligence

Informed Search



Instructor: Marco Alvarez
University of Rhode Island

[These slides were created by Dan Klein and Pieter Abbeel for CS188 Intro to AI at UC Berkeley (ai.berkeley.edu)]

2

Today


- Informed Search
 - Heuristics
 - Greedy Search
 - A* Search



3

Recap: Search

- Search problem:
 - States (configurations of the world)
 - Actions and costs
 - Successor function (world dynamics)
 - Start state and goal test
- Search tree:
 - Nodes: represent plans for reaching states
 - Plans have costs (sum of action costs)
- Search algorithm:
 - Systematically builds a search tree
 - Chooses an ordering of the fringe (unexplored nodes)
 - Optimal: finds least-cost plans



4



Uninformed Search



5

Uniform Cost Search


- Strategy: expand lowest path cost
- The good: UCS is complete and optimal!
- The bad:
 - Explores options in every "direction"
 - No information about goal location



[Demo: contours UCS empty (L3D1)]
[Demo: contours UCS pacman small maze (L3D3)]


6

Video of Demo Contours UCS Empty




7

Video of Demo Contours UCS Pacman Small Maze



8



Informed Search



9


Search Heuristics

- A heuristic is:
 - A function that *estimates* how close a state is to a goal
 - Designed for a particular search problem
 - Examples: Manhattan distance, Euclidean distance for pathing



10

Example: Heuristic Function




11

Effect of heuristics

Guide search *towards the goal* instead of *all over the place*



Informed



Uninformed

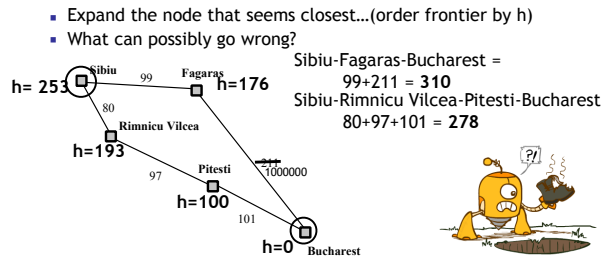
12

Greedy Search



13


Greedy Search



14

Greedy Search


- Strategy: expand a node that *seems* closest to a goal state
 - Heuristic: estimate of distance to nearest goal for each state
- A common case:
 - Best-first takes you straight to the (wrong) goal
- Worst-case: like a badly-guided DFS



[Demo: contours greedy empty (L3D1)]
[Demo: contours greedy pacman small maze (L3D4)]


15

Video of Demo Contours Greedy (Empty)



16

Video of Demo Contours Greedy (Pacman Small Maze)




17

A* Search




18

A* Search



19


Combining UCS and Greedy



20

When should A* terminate?


- Should we stop when we enqueue a goal?



- No: only stop when we dequeue a goal

21


Is A* Optimal?



- What went wrong?
- Actual bad goal cost < estimated good goal cost
- We need estimates to be less than actual costs!


22

Admissible Heuristics



23

Idea: Admissibility



Inadmissible (pessimistic) heuristics break optimality by trapping good plans on the fringe

Admissible (optimistic) heuristics slow down bad plans but never outweigh true costs

24

Admissible Heuristics

- A heuristic h is *admissible* (optimistic) if:

$$0 \leq h(n) \leq h^*(n)$$

where $h^*(n)$ is the true cost to a nearest goal

- Example:



- Coming up with admissible heuristics is most of what's involved in using A* in practice.

25


Properties of A*

Properties of A*

Uniform-Cost



A*




26

27

UCS vs A* Contours

- Uniform-cost expands equally in all “directions”




- A* expands mainly toward the goal, but does hedge its bets to ensure optimality



[Demo: contours UCS / greedy / A* empty (L3D1)]
[Demo: contours A* pacman small maze (L3D5)]


28

Video of Demo Contours (Empty) -- UCS




29

Video of Demo Contours (Empty) -- Greedy




30

Video of Demo Contours (Empty) - A*



31

Video of Demo Contours (Pacman Small Maze) - A*



32

Comparison




Greedy

Uniform Cost

A*

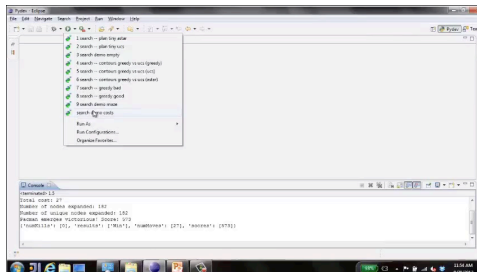
33

Video of Demo Pacman (Tiny Maze) - UCS / A*




34

Video of Demo Empty Water Shallow/Deep - Guess Algorithm



35

Creating Heuristics



36

Creating Admissible Heuristics


- Most of the work in solving hard search problems optimally is in coming up with admissible heuristics
- Often, admissible heuristics are solutions to *relaxed problems*, where new actions are available



- Inadmissible heuristics are often useful too

37

Example: 8 Puzzle





- What are the states?
- How many states?
- What are the actions?
- How many successors from the start state?
- What should the costs be?

38

8 Puzzle I

- Heuristic: Number of tiles misplaced
- Why is it admissible?
- $h(\text{start}) = 8$
- This is a *relaxed-problem* heuristic




Average nodes expanded when the optimal path has...			
	...4 steps	...8 steps	...12 steps
UCS	112	6,300	3.6×10^6
TILES	13	39	227

Statistics from Andrew Moore

39

8 Puzzle II

- What if we had an easier 8-puzzle where any tile could slide any direction at any time, ignoring other tiles?



- Total *Manhattan* distance

- Why is it admissible?

- $h(\text{start}) = 3 + 1 + 2 + \dots = 18$


Average nodes expanded when the optimal path has...			
	...4 steps	...8 steps	...12 steps
TILES	13	39	227
MANHATTAN	12	25	73

40

8 Puzzle III

- How about using the *actual cost* as a heuristic?

- Would it be admissible?
- Would we save on nodes expanded?
- What's wrong with it?



- With A^* : a trade-off between quality of estimate and work per node
 - As heuristics get closer to the true cost, you will expand fewer nodes but usually do more work per node to compute the heuristic itself

41