

AlphaGO

- Tic-Tac-Toe — PhD project @ Cambridge 1952
- Chess — Deep Blue 1997
- Jeopardy — Watson 2011
- Atari Games from raw pixels (Google) — 2014
- AlphaGo — Deepmind — 2016
 - GO game search space?
 - more than a GooGol larger than Chess!!
 - greater than #atoms in universe

1

[illegible]

2


AlphaGo

- Monte Carlo Tree Search
 - **Policy Network** – suggest moves
 - **Value Network** – estimates eventual winner
- Deep Neural Networks
 - 12 layers (supervised and reinforcement learning)
 - millions of connections
 - trained with 30 million of moves (human experts)
 - predicted 57%
 - networks by itself defeat state-of-the-art programs (based on trees)


3

Deep Learning Chip

- “On our chip we bring the data as close as possible to the processing units, and move the data as little as possible,”
- “Whereas many of the cores in a GPU share a single, large memory bank, each of the Eyeriss cores has its own memory”



Eyeriss has 168 processing elements (PE), each with its own memory.



MIT News

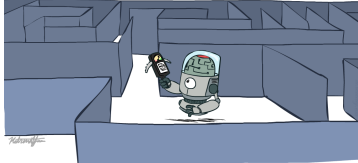
Advanced circuit enables mobile devices to perform "neural networks" modeled on the human brain.

Energy-friendly chip can perform powerful artificial-intelligence tasks

4

CS 188: Artificial Intelligence

A* Analysis




Instructor: Marco Alvarez
University of Rhode Island

[These slides were created by Dan Klein and Pieter Abbeel for CS188 Intro to AI at UC Berkeley (ai.berkeley.edu)]

5

Today


- Informed Search
 - A* Search (Analysis)
- Graph Search



A purple robot with a green eye and a 'GT' logo on its chest is holding a blue envelope with a yellow arrow pointing to it and a brown boot with steam coming out of it.

6

A* Search

A cartoon illustration of a mountain landscape. Two brown mountains with green grassy peaks are shown. A white path leads from the bottom center towards a cave entrance in the left mountain. A small yellow character with a red bow is running along the path towards the cave. Various items are scattered on the ground: a red diamond, a blue diamond, and a yellow gem. The sky is light blue with a few white clouds.

7

Combining UCS and Greedy

- Uniform-cost orders by path cost, or *backward cost* $g(n)$
- Greedy orders by goal proximity, or *forward cost* $h(n)$

Diagram illustrating a search graph with nodes S, a, b, c, d, e, and G. Edges and costs: S to a (1), a to b (1), b to c (1), a to d (3), d to e (1), e to G (2). Heuristic values: $h(S)=6$, $h(a)=5$, $h(b)=6$, $h(c)=7$, $h(d)=2$, $h(e)=1$, $h(G)=0$. A curved edge from S to e is labeled with cost 8.

Diagram illustrating a search tree with root a. Children: b, d, e. b has children c and G. d has child G. e has children d and G. Heuristic values: $g(a)=0$, $h(a)=6$; $g(b)=1$, $h(b)=5$; $g(c)=2$, $h(c)=6$; $g(d)=4$, $h(d)=2$; $g(e)=9$, $h(e)=1$; $g(c)=3$, $h(c)=7$; $g(G)=6$, $h(G)=0$; $g(d)=10$, $h(d)=2$; $g(G)=12$, $h(G)=0$.

- A* Search orders by the sum: $f(n) = g(n) + h(n)$

Example: Teg Grenager

8

```

graph LR
    S((S)) -- 1 --> A((A))
    A -- 3 --> G((G))
    S -- 5 --> G
    S_h["h = 7"] --- S
    A_h["h = 6"] --- A
    G_h["h = 0"] --- G

```

- What went wrong?
- Actual bad goal cost < estimated good goal cost
- We need estimates to be less than actual costs!

9

Admissible Heuristics

- A heuristic h is *admissible* (optimistic) if:

$$0 \leq h(n) \leq h^*(n)$$

where $h^*(n)$ is the true cost to a nearest goal


- Example:



- Coming up with admissible heuristics is most of what's involved in using A* in practice.

10

Optimality of A* Tree Search



11


Optimality of A* Tree Search

Assume:

- A is an optimal goal node
- B is a suboptimal goal node
- h is admissible

Claim:

- A will exit the fringe before B




12

Optimality of A* Tree Search: Blocking

Proof:

- Imagine B is on the fringe
- Some ancestor n of A is on the fringe, too (maybe A!)
- Claim: n will be expanded before B

- $f(n)$ is less or equal to $f(A)$



$$\begin{aligned} f(n) &= g(n) + h(n) \\ f(n) &\leq g(A) \\ g(A) &= f(A) \end{aligned}$$

Definition of f-cost
Admissibility of h
 $h = 0$ at a goal


13

Optimality of A* Tree Search: Blocking

Proof:

- Imagine B is on the fringe
- Some ancestor n of A is on the fringe, too (maybe A!)
- Claim: n will be expanded before B

- $f(n)$ is less or equal to $f(A)$
- $f(A)$ is less than $f(B)$



$$\begin{aligned} g(A) &< g(B) && B \text{ is suboptimal} \\ f(A) &< f(B) && h = 0 \text{ at a goal} \end{aligned}$$

14


Optimality of A* Tree Search: Blocking

Proof:

- Imagine B is on the fringe
- Some ancestor n of A is on the fringe, too (maybe A!)
- Claim: n will be expanded before B

- $f(n)$ is less or equal to $f(A)$
- $f(A)$ is less than $f(B)$
- n expands before B

- All ancestors of A expand before B
- A expands before B
- A* search is optimal



$$f(n) \leq f(A) < f(B)$$

15

Semi-Lattice of Heuristics

16

Trivial Heuristics, Dominance

- Dominance: $h_a \geq h_c$ if

$$\forall n : h_a(n) \geq h_c(n)$$


- Heuristics form a semi-lattice:

- Max of admissible heuristics is admissible

$$h(n) = \max(h_a(n), h_b(n))$$


- Trivial heuristics

- Bottom of lattice is the zero heuristic (what does this give us?)
- Top of lattice is the exact heuristic



17


Graph Search



18

Tree Search: Extra Work!


- Failure to detect repeated states can cause exponentially more work.



19

Graph Search

- In BFS, for example, we shouldn't bother expanding the circled nodes (why?)



20


Graph Search

- Idea: never expand a state twice
- How to implement:
 - Tree search + set of expanded states ("closed set")
 - Expand the search tree node-by-node, but...
 - Before expanding a node, check to make sure its state has never been expanded before
 - If not new, skip it, if new add to closed set
- Important: store the closed set as a set, not a list
- Can graph search wreck completeness? Why/why not?
- How about optimality?

states (Graph)
nodes (Tree)


21

A* Graph Search Gone Wrong?




22

Consistency of Heuristics



23

Optimality of A* Graph Search




24

Optimality of A* Graph Search

- Sketch: consider what A* does with a consistent heuristic:


- Fact 1: In tree search, A* expands nodes in increasing total f value (f-contours)
- Fact 2: For every state s, nodes that reach s optimally are expanded before nodes that reach s suboptimally
- Result: A* graph search is optimal



25

Optimality

- Tree search:
 - A* is optimal if heuristic is admissible
 - UCS is a special case (h = 0)
- Graph search:
 - A* optimal if heuristic is consistent
 - UCS optimal (h = 0 is consistent)
- Consistency implies admissibility
- In general, most natural admissible heuristics tend to be consistent, especially if from relaxed problems



26

A*: Summary



27

A*: Summary

- A* uses both backward costs and (estimates of) forward costs
- A* is optimal with admissible / consistent heuristics
- Heuristic design is key: often use relaxed problems



28

Tree Search Pseudo-Code

```
function TREE-SEARCH(problem, fringe) return a solution, or failure
  fringe ← INSERT(MAKE-NODE(INITIAL-STATE[problem]), fringe)
  loop do
    if fringe is empty then return failure
    node ← REMOVE-FRONT(fringe)
    if GOAL-TEST(problem, STATE[node]) then return node
    for child-node in EXPAND(STATE[node], problem) do
      fringe ← INSERT(child-node, fringe)
  end
end
```

29

Graph Search Pseudo-Code


```
function GRAPH-SEARCH(problem, fringe) return a solution, or failure
  closed ← an empty set
  fringe ← INSERT(MAKE-NODE(INITIAL-STATE[problem]), fringe)
  loop do
    if fringe is empty then return failure
    node ← REMOVE-FRONT(fringe)
    if GOAL-TEST(problem, STATE[node]) then return node
    if STATE[node] is not in closed then
      add STATE[node] to closed
      for child-node in EXPAND(STATE[node], problem) do
        fringe ← INSERT(child-node, fringe)
      end
    end
  end
```

30

Optimality of A* Graph Search

- Consider what A* does:
 - Expands nodes in increasing total f value (f -contours)
- Reminder: $f(n) = g(n) + h(n) = \text{cost to } n + \text{heuristic}$
- Proof idea: the optimal goal(s) have the lowest f value, so it must get expanded first

There's a problem with this argument. What are we assuming is true?




31

Optimality of A* Graph Search

Proof:

- New possible problem: some n on path to G^* isn't in queue when we need it, because some worse n' for the same state dequeued and expanded first (disaster!)
- Take the highest such n in tree
- Let p be the ancestor of n that was on the queue when n' was popped
- $f(p) < f(n)$ because of consistency
- $f(n) < f(n')$ because n' is suboptimal
- p would have been expanded before n'
- Contradiction!



32

A* Applications

- Video games
- Pathing / routing problems
- Resource planning problems
- Robot motion planning
- Language analysis
- Machine translation
- Speech recognition
- ...



33