

## Announcements

- Optional Mini-Contest (search)
  - due Friday 3/1 at 11:59p
  - Use Autolab for submissions
- Homework 3: Games
  - Has been released, due 3/3 at 11:59p.
  - Few questions about Utilities (next lecture)
- Project 1
  - Please resubmit on Autolab today
  - single ZIP file (search.py and searchAgents.py only – no folders)

**New Rules:**

**Homework:**  
submit answers on edX  
individual work/individual submission

**Projects:**  
submit source code on Autolab  
team work/individual submission


**Contest:**  
submit source code on Autolab  
team work/team submission

All assignments will be due at 11:59pm EST

No late submissions  
No submissions by e-mail

1

## Project 1 Mini-Contest



2

## CS 188: Artificial Intelligence Stochastic Games




Instructor: Marco Alvarez  
University of Rhode Island

[These slides were created by Dan Klein and Pieter Abbeel for CS188 Intro to AI at UC Berkeley. All CS188 materials are available at <http://ai.berkeley.edu>.]


3

## Recap - Adversarial Games



4

## Minimax Example



5

## Minimax Implementation

```
def max-value(state):
    initialize v = -∞
    for each successor of state:
        v = max(v, min-value(successor))
    return v
```

```
def min-value(state):
    initialize v = +∞
    for each successor of state:
        v = min(v, max-value(successor))
    return v
```

$$V(s) = \max_{s' \in \text{Successors}(s)} V(s')$$

$$V(s') = \min_{s \in \text{Successors}(s')} V(s)$$

6

## Minimax Implementation (Dispatch)

```
def value(state):
    if the state is a terminal state: return the state's utility
    if the next agent is MAX: return max-value(state)
    if the next agent is MIN: return min-value(state)
```


```
def max-value(state):
    initialize v = -∞
    for each successor of state:
        v = max(v, value(successor))
    return v
```

```
def min-value(state):
    initialize v = +∞
    for each successor of state:
        v = min(v, value(successor))
    return v
```

7


## Minimax Efficiency

- How efficient is minimax?
  - Just like (exhaustive) DFS
  - Time:  $O(b^m)$
  - Space:  $O(bm)$
- Example: For chess,  $b \approx 35$ ,  $m \approx 100$ 
  - Exact solution is completely infeasible
  - But, do we need to explore the whole tree?



8


## Evaluation Functions



9

## Evaluation Functions

- Evaluation functions score non-terminals in depth-limited search



- Ideal function: returns the actual minimax value of the position
- In practice: typically weighted linear sum of features:

$$Eval(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$$

- e.g.  $f_1(s) = (\text{num white queens} - \text{num black queens})$ , etc.

10


## Resource Limits

- Problem: In realistic games, cannot search to leaves!

- Solution: Depth-limited search
  - Instead, search only to a limited depth in the tree
  - Replace terminal utilities with an evaluation function for non-terminal positions


- Example:
  - Suppose we have 100 seconds, can explore 10K nodes / sec
  - So can check 1M nodes per move
  - $\alpha$ - $\beta$  reaches about depth 8 - decent chess program

- Guarantee of optimal play is gone
- More plies makes a BIG difference
- Use iterative deepening for an anytime algorithm



11


## Why Pacman Starves



- A danger of replanning agents!
  - He knows his score will go up by eating the dot now (west, east)
  - He knows his score will go up just as much by eating the dot later (east, west)
  - There are no point-scoring opportunities after eating the dot (within the horizon, two here)
  - Therefore, waiting seems just as good as eating: he may go east, then back west in the next round of replanning!


12

## Game Tree Pruning



13

## Minimax Pruning



14

## Alpha-Beta Implementation

$\alpha$ : MAX's best option on path to root  
 $\beta$ : MIN's best option on path to root

```
def max-value(state,  $\alpha$ ,  $\beta$ ):
    initialize  $v = -\infty$ 
    for each successor of state:
         $v = \max(v, \text{value}(\text{successor}, \alpha, \beta))$ 
        if  $v \geq \beta$  return  $v$ 
         $\alpha = \max(\alpha, v)$ 
    return  $v$ 
```

```
def min-value(state,  $\alpha$ ,  $\beta$ ):
    initialize  $v = +\infty$ 
    for each successor of state:
         $v = \min(v, \text{value}(\text{successor}, \alpha, \beta))$ 
        if  $v \leq \alpha$  return  $v$ 
         $\beta = \min(\beta, v)$ 
    return  $v$ 
```

15

## Alpha-Beta Pruning Properties


- This pruning has no effect on minimax value computed for the root!

- Values of intermediate nodes might be wrong
  - Important: children of the root may have the wrong value
  - So the most naive version won't let you do action selection

- Good child ordering improves effectiveness of pruning


- With "perfect ordering":
  - Time complexity drops to  $O(b^{n/2})$
  - Doubles solvable depth!
  - Full search of, e.g. chess, is still hopeless...

- This is a simple example of metareasoning (computing about what to compute)




16

## Uncertain Outcomes



17

## Worst-Case vs. Average Case




Idea: Uncertain outcomes controlled by chance, not an adversary!

18

## Expectimax Search

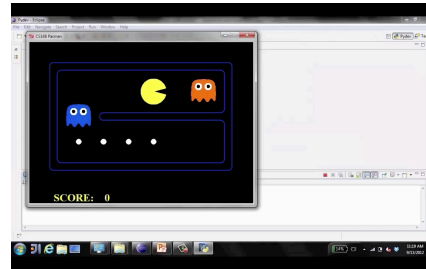
- Why wouldn't we know what the result of an action will be?
  - Explicit randomness: rolling dice
  - Unpredictable opponents: the ghosts respond randomly
  - Actions can fail: when moving a robot, wheels might slip
- Values should now reflect average-case (expectimax) outcomes, not worst-case (minimax) outcomes
- Expectimax search: compute the average score under optimal play
  - Max nodes as in minimax search
  - Chance nodes are like min nodes but the outcome is uncertain
  - Calculate their expected utilities
  - i.e. take weighted average (expectation) of children
- Later, we'll learn how to formalize the underlying uncertain-result problems as Markov Decision Processes



[Demo: min vs exp (L7D1.2)]


19

## Video of Demo Minimax vs Expectimax (Min)



20

## Video of Demo Minimax vs Expectimax (Exp)



21

## Expectimax Pseudocode

```
def value(state):
    if the state is a terminal state: return the state's utility
    if the next agent is MAX: return max-value(state)
    if the next agent is EXP: return exp-value(state)
```


```
def max-value(state):
    initialize v = -∞
    for each successor of state:
        v = max(v, value(successor))
    return v
```

```
def exp-value(state):
    initialize v = 0
    for each successor of state:
        p = probability(successor)
        v += p * value(successor)
    return v
```

22

## Expectimax Pseudocode


```
def exp-value(state):
    initialize v = 0
    for each successor of state:
        p = probability(successor)
        v += p * value(successor)
    return v
```



$$v = (1/2) (8) + (1/3) (24) + (1/6) (-12) = 10$$


23

## Expectimax Example




24

## Expectimax Pruning?



25

## Depth-Limited Expectimax



26

## Probabilities



27

## Reminder: Probabilities

- A random variable represents an event whose outcome is unknown
- A probability distribution is an assignment of weights to outcomes

### Example: Traffic on freeway


- Random variable:  $T$  = whether there's traffic
- Outcomes:  $T$  in {none, light, heavy}
- Distribution:  $P(T=\text{none}) = 0.25$ ,  $P(T=\text{light}) = 0.50$ ,  $P(T=\text{heavy}) = 0.25$

### Some laws of probability (more later):

- Probabilities are always non-negative
- Probabilities over all possible outcomes sum to one

### As we get more evidence, probabilities may change:

- $P(T=\text{heavy}) = 0.25$ ,  $P(T=\text{heavy} \mid \text{Hour}=8\text{am}) = 0.60$
- We'll talk about methods for reasoning and updating probabilities later




28

## Reminder: Expectations

- The expected value of a function of a random variable is the average, weighted by the probability distribution over outcomes


### Example: How long to get to the airport?



29

## What Probabilities to Use?


- In expectimax search, we have a probabilistic model of how the opponent (or environment) will behave in any state
  - Model could be a simple uniform distribution (roll a die)
  - Model could be sophisticated and require a great deal of computation
  - We have a chance node for any outcome out of our control: opponent or environment
  - The model might say that adversarial actions are likely!
- For now, assume each chance node magically comes along with probabilities that specify the distribution over its outcomes



30

## Quiz: Informed Probabilities

- Let's say you know that your opponent is actually running a depth 2 minimax, using the result 80% of the time, and moving randomly otherwise
- Question: What tree search should you use?




### Answer: Expectimax!

- To figure out EACH chance node's probabilities, you have to run a simulation of your opponent
- This kind of thing gets very slow very quickly
- Even worse if you have to simulate your opponent simulating you...
- ...except for minimax, which has the nice property that it all collapses into one game tree

31

## Modeling Assumptions



32

## The Dangers of Optimism and Pessimism

**Dangerous Optimism**  
Assuming chance when the world is adversarial




**Dangerous Pessimism**  
Assuming the worst case when it's not likely



33

## Assumptions vs. Reality



|                   | Adversarial Ghost           | Random Ghost               |
|-------------------|-----------------------------|----------------------------|
| Minimax Pacman    | Won 5/5<br>Avg. Score: 483  | Won 5/5<br>Avg. Score: 493 |
| Expectimax Pacman | Won 1/5<br>Avg. Score: -303 | Won 5/5<br>Avg. Score: 503 |


Results from playing 5 games

Pacman used depth 4 search with an eval function that avoids trouble  
Ghost used depth 2 search with an eval function that seeks Pacman

[Demos: world assumptions (L7D3,4,5,6)]


34

## Video of Demo World Assumptions Random Ghost - Expectimax Pacman



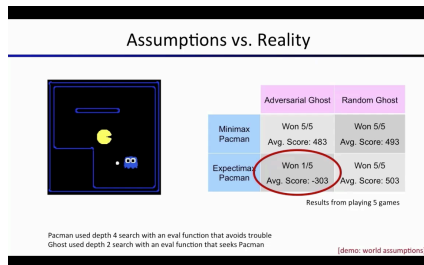
35

## Video of Demo World Assumptions Adversarial Ghost - Minimax Pacman




36

## Video of Demo World Assumptions Adversarial Ghost - Expectimax Pacman




37

## Video of Demo World Assumptions Random Ghost - Minimax Pacman



38


## Other Game Types



39

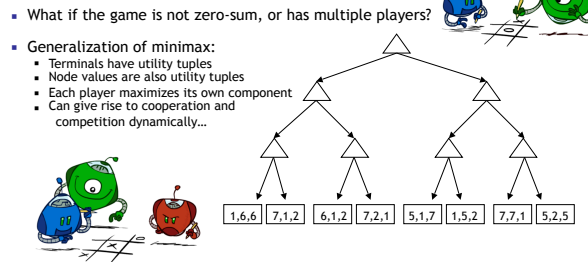
## Mixed Layer Types

- E.g. Backgammon
- Expectiminimax
  - Environment is an extra "random agent" player that moves after each min/max agent
  - Each node computes the appropriate combination of its children



40

## Multi-Agent Utilities



41