

# Lecture 2: Image Classification pipeline

**Image Classification:** a core task in Computer Vision



(assume given set of discrete labels)  
{dog, cat, truck, plane, ...}


→ cat

**The problem:**  
*semantic gap*

Images are represented as  
3D arrays of numbers, with  
integers between [0, 255].

E.g.  
300 x 100 x 3

(3 for 3 color channels RGB)




Fei-Fei Li & Andrej Karpathy & Justin Johnson Lecture 2 - 1 6 Jan 2016

Fei-Fei Li & Andrej Karpathy & Justin Johnson Lecture 2 - 6 6 Jan 2016

Fei-Fei Li & Andrej Karpathy & Justin Johnson Lecture 2 - 7 6 Jan 2016


Challenges: Viewpoint Variation



Challenges: Illumination



Challenges: Deformation




Fei-Fei Li & Andrej Karpathy & Justin Johnson Lecture 2 - 8 6 Jan 2016


Fei-Fei Li & Andrej Karpathy & Justin Johnson Lecture 2 - 9 6 Jan 2016

Fei-Fei Li & Andrej Karpathy & Justin Johnson Lecture 2 - 10 6 Jan 2016

Challenges: Occlusion



Challenges: Background clutter



Challenges: Intra-class variation



Fei-Fei Li & Andrej Karpathy & Justin Johnson Lecture 2 - 11 6 Jan 2016

Fei-Fei Li & Andrej Karpathy & Justin Johnson Lecture 2 - 12 6 Jan 2016

Fei-Fei Li & Andrej Karpathy & Justin Johnson Lecture 2 - 13 6 Jan 2016


## An image classifier

```
def predict(image):
    # ???
    return class_label
```

Unlike e.g. sorting a list of numbers,

**no obvious way** to hard-code the algorithm for recognizing a cat, or other classes.

Attempts have been made



## Data-driven approach:

1. Collect a dataset of images and labels
2. Use Machine Learning to train an image classifier
3. Evaluate the classifier on a withheld set of test images

```
def train(train_images, train_labels):
    # build a model for images -> labels...
    return model

def predict(model, test_images):
    # predict test_labels using the model...
    return test_labels
```

Example training set



Fei-Fei Li & Andrej Karpathy & Justin Johnson Lecture 2 - 14 6 Jan 2016

Fei-Fei Li & Andrej Karpathy & Justin Johnson Lecture 2 - 15 6 Jan 2016

Fei-Fei Li & Andrej Karpathy & Justin Johnson Lecture 2 - 16 6 Jan 2016

## First classifier: Nearest Neighbor Classifier

```
def train(train_images, train_labels):
    # build a model for images -> labels...
    return model

def predict(model, test_images):
    # predict test_labels using the model...
    return test_labels
```

Remember all training images and their labels

Predict the label of the most similar training image

Example dataset: CIFAR-10  
10 labels  
50,000 training images, each image is tiny: 32x32  
10,000 test images.



Fei-Fei Li & Andrej Karpathy & Justin Johnson Lecture 2 - 17 6 Jan 2016

Fei-Fei Li & Andrej Karpathy & Justin Johnson Lecture 2 - 18 6 Jan 2016

Example dataset: CIFAR-10  
10 labels  
50,000 training images  
10,000 test images.



For every test image (first column), examples of nearest neighbors in rows



Fei-Fei Li & Andrej Karpathy & Justin Johnson Lecture 2 - 19 6 Jan 2016

How do we compare the images? What is the **distance metric**?

**L1 distance:**  $d_1(I_1, I_2) = \sum_p |I_1^p - I_2^p|$

test image	training image	pixel-wise absolute value differences
56 32 10 18	10 20 24 17	46 12 14 1
90 23 128 133	8 10 89 100	82 13 39 33
24 26 178 200	12 16 178 170	12 10 0 30
2 0 255 220	4 32 233 112	2 32 22 108

=

add → 456

```
import numpy as np

class NearestNeighbor:
    def __init__(self):
        pass

    def train(self, X, y):
        """ X is N x D where each row is an example. Y is 1-dimension of size N """
        # the nearest neighbor classifier simply remembers all the training data
        self.Xtr = X
        self.ytr = y

    def predict(self, X):
        """ X is N x D where each row is an example we wish to predict label for """
        num_test = X.shape[0]
        # lets make sure that the output type matches the input type
        Ypred = np.zeros(num_test, dtype = self.ytr.dtype)

        # loop over all test rows
        for i in xrange(num_test):
            # find the nearest training image to the i'th test image
            # using the L1 distance (sum of absolute value differences)
            distances = np.sum(np.abs(self.Xtr - X[i,:]), axis = 1)
            min_index = np.argmin(distances) # get the index with smallest distance
            Ypred[i] = self.ytr[min_index] # predict the label of the nearest example

        return Ypred
```

Nearest Neighbor classifier

Fei-Fei Li & Andrej Karpathy & Justin Johnson Lecture 2 - 20 6 Jan 2016

Fei-Fei Li & Andrej Karpathy & Justin Johnson Lecture 2 - 21 6 Jan 2016

```
import numpy as np

class NearestNeighbor:
    def __init__(self):
        pass

    def train(self, X, y):
        """ X is N x D where each row is an example. Y is 1-dimension of size N """
        # the nearest neighbor classifier simply remembers all the training data
        self.Xtr = X
        self.ytr = y

    def predict(self, X):
        """ X is N x D where each row is an example we wish to predict label for """
        num_test = X.shape[0]
        # lets make sure that the output type matches the input type
        Ypred = np.zeros(num_test, dtype = self.ytr.dtype)

        # loop over all test rows
        for i in xrange(num_test):
            # find the nearest training image to the i'th test image
            # using the L1 distance (sum of absolute value differences)
            distances = np.sum(np.abs(self.Xtr - X[i,:]), axis = 1)
            min_index = np.argmin(distances) # get the index with smallest distance
            Ypred[i] = self.ytr[min_index] # predict the label of the nearest example

        return Ypred
```

Nearest Neighbor classifier

remember the training data

Fei-Fei Li & Andrej Karpathy & Justin Johnson Lecture 2 - 22 6 Jan 2016

Nearest Neighbor classifier

- for every test image:
  - find nearest train image with L1 distance
  - predict the label of nearest training image

Nearest Neighbor classifier

Fei-Fei Li &amp; Andrej Karpathy &amp; Justin Johnson      Lecture 2 - 24      6 Jan 2016

Nearest Neighbor classifier











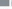






















































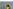






















































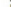









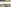
Fei-Fei Li &amp; Andrej Karpathy &amp; Justin Johnson      Lecture 2 - 25      6 Jan 2016

[illegible]
$$d_1(I_1, I_2) = \sum_p |I_1^p - I_2^p|$$
$$d_2(I_1, I_2) = \sqrt{\sum_p (I_1^p - I_2^p)^2}$$

Fei-Fei Li &amp; Andrej Karpathy &amp; Justin Johnson      Lecture 2 - 27      6 Jan 2016

The figure consists of three side-by-side scatter plots. The first plot, titled 'the data', shows a distribution of red, green, and blue points. The second plot, titled '1-NN classifier', shows the same points with decision boundaries that are highly irregular and complex, following the local distribution of the training points. The third plot, titled '5-NN classifier', shows the same points with much smoother and more regular decision boundaries, indicating better generalization.

Fei-Fei Li &amp; Andrej Karpathy &amp; Justin Johnson      Lecture 2 - 28      6 Jan 2016

airplane																	
automobile																	
bird																	
cat																	
deer																	
dog																	
frog																	
horse																	

Fei-Fei Li &amp; Andrej Karpathy &amp; Justin Johnson      Lecture 2 - 30      6 Jan 2016

Fei-Fei Li &amp; Andrej Karpathy &amp; Justin Johnson      Lecture 2 - 30      6 Jan 2016

Fei-Fei Li &amp; Andrej Karpathy &amp; Justin Johnson      Lecture 2 - 31      6 Jan 2016

What is the best **distance** to use?  
What is the best value of **k** to use?


i.e. how do we set the **hyperparameters**?

What is the best **distance** to use?  
What is the best value of **k** to use?




i.e. how do we set the **hyperparameters**?

Very problem-dependent.  
Must try them all out and see what works best.

Try out what hyperparameters work best on test set.



Trying out what hyperparameters work best on test set:  
Very bad idea. The test set is a proxy for the generalization performance!  
Use only **VERY SPARINGLY**, at the end.



Example of  
5-fold cross-validation  
for the value of **k**.


Each point: single  
outcome.

The line goes  
through the mean,  
bars indicated  
standard  
deviation

(Seems that  $k \sim 7$  works best  
for this data)

k-Nearest Neighbor on images **never used**.

- terrible performance at test time
- distance metrics on level of whole images can be very unintuitive




(all 3 images have same L2 distance to the one on the left)

## Summary


- **Image Classification:** We are given a **Training Set** of labeled images, asked to predict labels on **Test Set**. Common to report the **Accuracy** of predictions (fraction of correctly predicted images)
- We introduced the **k-Nearest Neighbor Classifier**, which predicts the labels based on nearest images in the training set
- We saw that the choice of distance and the value of **k** are **hyperparameters** that are tuned using a **validation set**, or through **cross-validation** if the size of the data is small.
- Once the best set of hyperparameters is chosen, the classifier is evaluated once on the test set, and reported as the performance of kNN on that data.




# Linear Classification




## Parametric approach




## Parametric approach: Linear classifier




## Parametric approach: Linear classifier




## Parametric approach: Linear classifier



## Example with an image with 4 pixels, and 3 classes (cat/dog/ship)




## Interpreting a Linear Classifier



$$f(x_i, W, b) = Wx_i + b$$

Q: what does the linear classifier do, in English?


## Interpreting a Linear Classifier



$$f(x_i, W, b) = Wx_i + b$$

Example trained weights of a linear classifier trained on CIFAR-10:

## Interpreting a Linear Classifier



## Interpreting a Linear Classifier



So far: We defined a (linear) **score function**:  $f(x_i, W, b) = Wx_i + b$

Example class scores for 3 images, with a random  $W$ :

airplane	-3.45	-0.51	3.42
automobile	-8.87	<b>6.04</b>	4.64
bird	0.09	5.31	2.65
cat	<b>2.9</b>	-4.22	5.1
deer	4.48	-4.19	2.64
dog	8.02	3.58	5.55
frog	3.78	4.49	<b>-4.34</b>
horse	1.06	-4.37	-1.5
ship	-0.36	-2.09	-4.79
truck	-0.72	-2.93	6.14

Coming up:

- Loss function (quantifying what it means to have a "good"  $W$ )
- Optimization (start with random  $W$  and find a  $W$  that minimizes the loss)
- ConvNets! (tweak the functional form of  $f$ )